

# WCTF 2018 binja Editorial

# rsWC

Author: @Charo\_IT

# Overview


- A simple memo manager which uses original heap allocator
- You are dropped into an unprivileged shell
- Can you exploit the binary and gain your privilege?

# About the original allocator

```
struct {  
    void *heap_base;  
    void *top;  
    size_t heap_size;  
    unsigned long num_chunks;  
    chunkinfo chunks[(0x1000 - 0x20) / sizeof(chunkinfo)];  
    // typedef struct {  
    //     void *ptr;  
    //     size_t size;  
    // } chunkinfo;  
} *arena;
```

# Initializing the heap

```
1. mmap(NULL, 0x1000, PROT_NONE, ...);  
   // map guard page
```



# Initializing the heap

2. arena =

```
mmap(NULL, 0x1000, PROT_READ | PROT_WRITE, ...);  
// map area for arena
```



# Initializing the heap

```
3. arena->heap_base =  
    mmap(NULL, heap_size, PROT_READ | PROT_WRITE,  
        ...); // map data area
```



# On malloc

1. Round up requested size to multiple of 0x10
2. Scan arena->chunks and find a chunk which satisfies  
`(chunk->size & 1) == 0` // chunk is not in use  
&& `aligned_requested_size <= arena->chunk[idx].size`
3. If found, set LSB of arena->chunks[idx].size and return  
arena->chunks[idx].ptr
4. If not found, create a new chunk at arena->top
5. Add new chunk to arena->chunks
6. Return the address of new chunk



## On free

1. Find target chunk from arena->chunks
2. Clear LSB of arena->chunks[idx].size

# Vulnerability

On malloc:

1. Round up requested size to multiple of 0x10
2. Scan arena->chunks and find a chunk which satisfies  
`(chunk->size & 1) == 0` // chunk is not in use  
&& `aligned_requested_size <= arena->chunk[idx].size`
3. If found, set LSB of arena->chunks[idx].size and return  
arena->chunks[idx].ptr
4. If not found, create a new chunk at arena->top
5. Add new chunk to arena->chunks
6. Return the address of new chunk

# Vulnerability


```
... // no chunks available so use top
... if(arena->heap_base + arena->heap_size < arena->top + size){
...     return NULL;
... }
... p = arena->top;
... arena->top += size;
... arena->chunks[arena->num_chunks].ptr = p;
... arena->chunks[arena->num_chunks].size = size | 1;
... arena->num_chunks++;

... return p;
}
```

allocating many chunks will make  
arena->chunks overflow

# Is this exploitable?

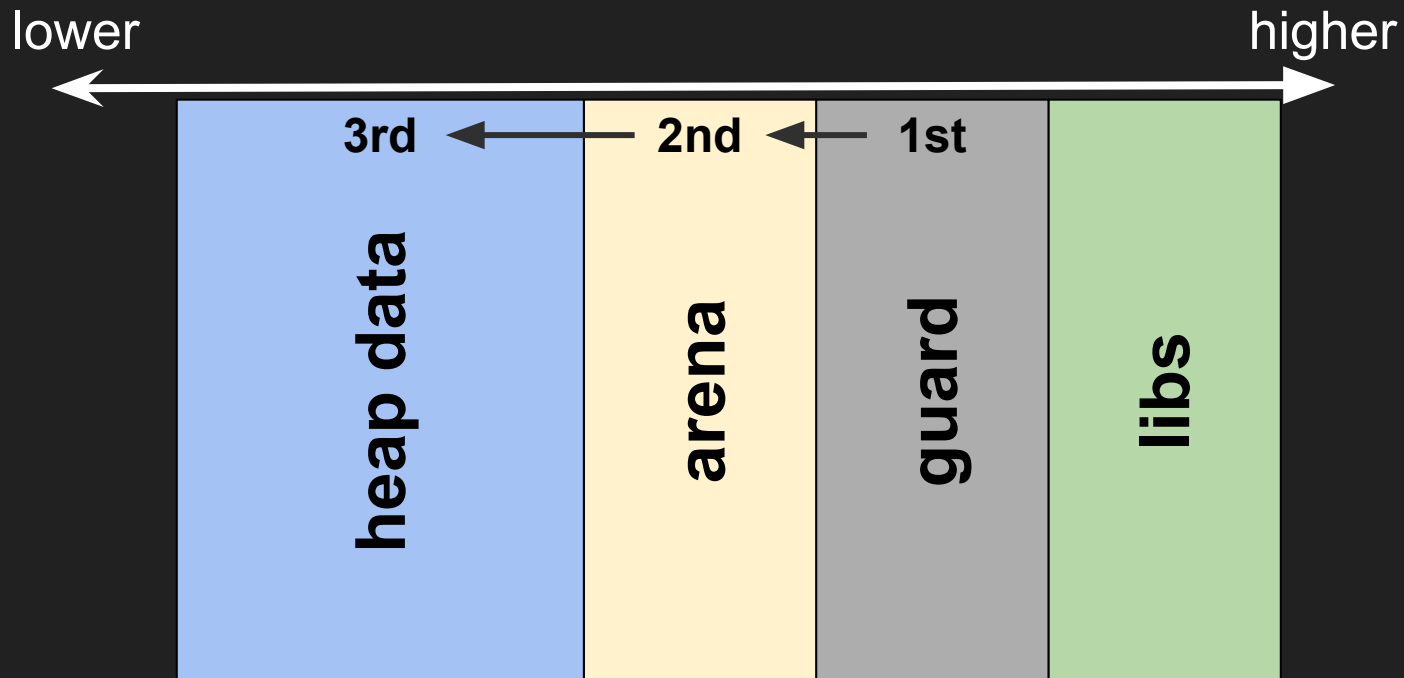
arena->chunks will overflow but we can't overwrite anything because of guard page :(



But wait...

The initialization process relies on the assumption that **mmap tries to allocate pages from higher to lower address**.

Can we break it?



# About mmap layout

```
void arch_pick_mmap_layout(struct mm_struct *mm)
{
→   unsigned long random_factor = 0UL;

→   if (current->flags & PF_RANDOMIZE)
→       random_factor = arch_mmap_rnd();

→   mm->mmap_legacy_base = TASK_UNMAPPED_BASE + random_factor;

→   if (mmap_is_legacy()) {
→       mm->mmap_base = mm->mmap_legacy_base;
→       mm->get_unmapped_area = arch_get_unmapped_area;
→   } else {
→       mm->mmap_base = mmap_base(random_factor);
→       mm->get_unmapped_area = arch_get_unmapped_area_topdown;
→   }
}
```

# About mmap layout

```
void arch_pick_mmap_layout(struct mm_struct *mm)
{
    unsigned long random_factor = 0ULL;

    if (curr
        rand

    mm->mmap_base = mmap_base(random_factor);

    if (mmap_is_legacy()) {
        mm->mmap_base = mm->mmap_legacy_base;
        mm->get_unmapped_area = arch_get_unmapped_area;
    } else {
        mm->mmap_base = mmap_base(random_factor);
        mm->get_unmapped_area = arch_get_unmapped_area_topdown;
    }
}
```

If some condition is met, use bottom-up  
(from lower to higher address) layout

# About mmap layout

```
void arch_pick_mmap_layout(struct mm_struct *mm)
{
    unsigned long random_factor = 0UL;

    if (current->flags & PF_RANDOMIZE)
        random_factor = (unsigned long)current->uid;

    mm->mmap_base = mmap_base(random_factor);
    mm->get_unmapped_area = arch_get_unmapped_area;

    if (mmap_base == 0)
    {
        mm->mmap_base = mmap_base(random_factor);
        mm->get_unmapped_area = arch_get_unmapped_area_topdown;
    }
}
```

If not, use top-down (from higher to lower address) layout



# About mmap layout

```
static int mmap_is_legacy(void)
{
→   if (current->personality & ADDR_COMPAT_LAYOUT)
→       return 1;

→   if (rlimit(RLIMIT_STACK) == RLIM_INFINITY)
→       return 1;

→   return sysctl_legacy_va_layout;
}
```

# Let's confirm it

```
$ ./test
```

```
[mmap 1] 0x7fc03cd9b000  
[mmap 2] 0x7fc03cd9a000  
[mmap 3] 0x7fc03cd99000  
[mmap 4] 0x7fc03cd98000  
[mmap 5] 0x7fc03cd97000
```



higher to lower

```
$ ulimit -s unlimited; ./test
```

```
[mmap 1] 0x2b0b265a1000  
[mmap 2] 0x2b0b265a2000  
[mmap 3] 0x2b0b265a3000  
[mmap 4] 0x2b0b265a4000  
[mmap 5] 0x2b0b265a5000
```




lower to higher

Note: This behavior has been removed in Linux 4.13  
(The challenge was running on Linux 4.4)

# Is this challenge exploitable?

We can change mmap layout to bottom-up style since we can do "ulimit -s unlimited".

So yes, it is exploitable!



# Solution

1. `ulimit -s unlimited`
2. Launch the binary
3. Allocate many chunks to make arena->chunks overflow
4. Break link list of memo
5. GOT leak & GOT overwrite (eg. `atoi` -> `gets`)
6. Hijack the control flow and read the flag file

# Truth

Author: @Charo\_IT

# Overview

- .NET reversing challenge which looks very easy at a first glance

# Static analysis with dnSpy

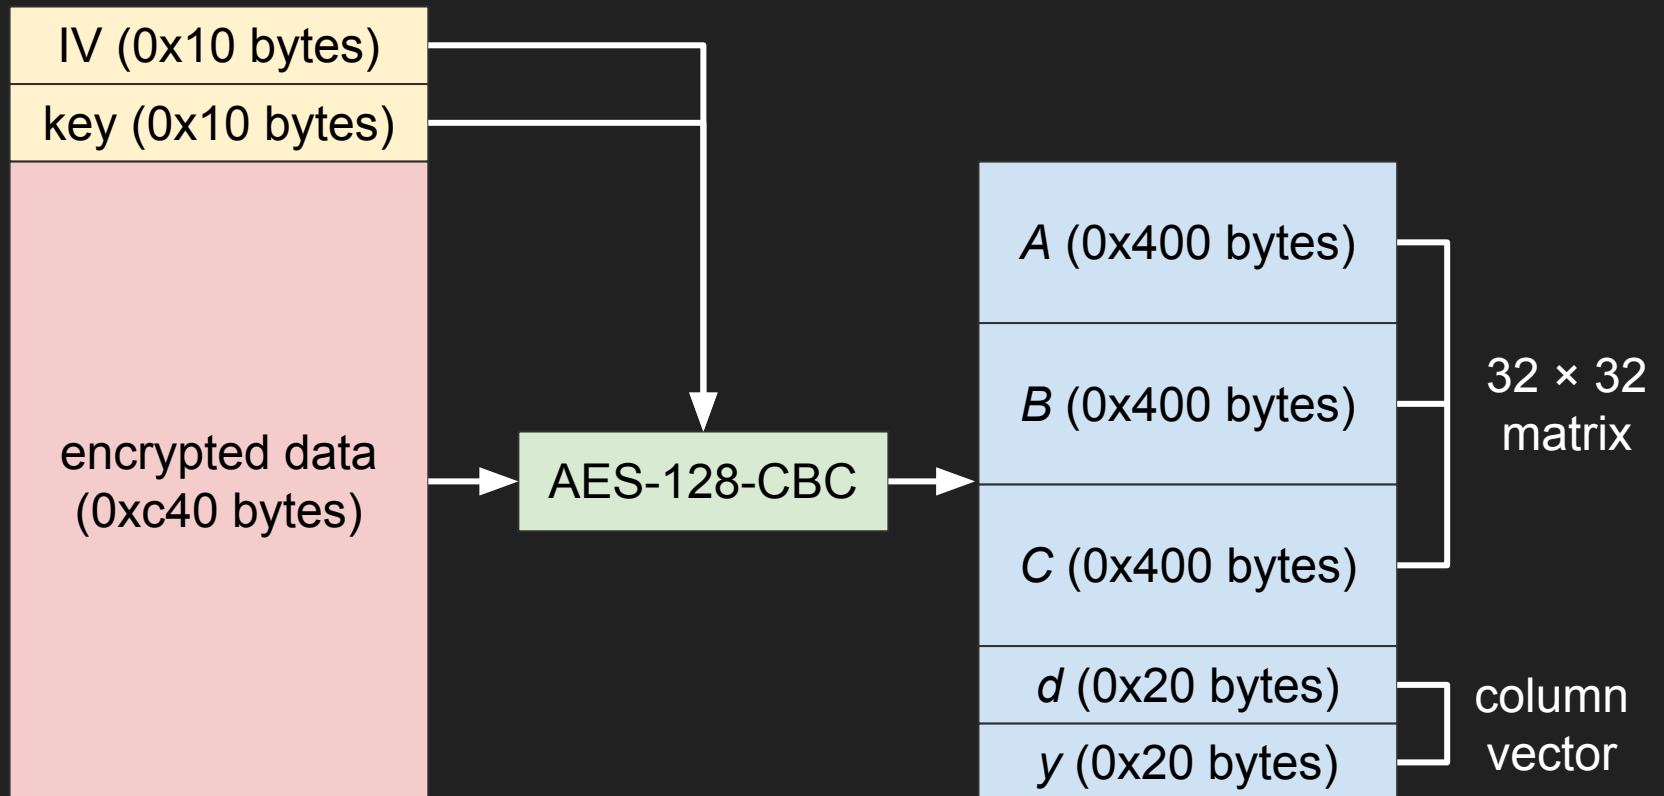
```
private static void Main() {  
→ Console.Write("Enter your flag: ");  
→ if (Lib.Verify(Console.ReadLine().Trim())) {  
→     → Console.WriteLine("Great :)");  
→     → return;  
→ }  
→ Console.WriteLine("Wrong :(");  
}
```

# Static analysis with dnSpy

```
public static bool Verify(string s) {  
→ byte[] bytes = Encoding.ASCII.GetBytes(s);  
→ if (bytes.Length != 32) {  
→ |→ return false;  
→ }  
→ byte[] array = Lib.Func2(); Read and decrypt resource  
→ Lib.Func3(array, bytes);  
→ Lib.Func4(array, bytes);  
→ Lib.Func5(array, bytes);  
→ for (int i = 0; i < 32; i++) {  
→ |→ if (bytes[i] != array[3104 + i]) {  
→ |→ |→ return false;  
→ |→ }  
→ }  
→ return true;  
}
```



# Deconstructing resource



# Static analysis with dnSpy

```
public static bool Verify(string s) {  
→ byte[] bytes = Encoding.ASCII.GetBytes(s);  $x_0 := \text{input}$   
→ if (bytes.Length != 32) {  
→ |→ return false;  
→ }  
→ byte[] array = Lib.Func2();  
→ Lib.Func3(array, bytes);  $\text{Func3: } x_1 := Ax_0$   
→ Lib.Func4(array, bytes);  $\text{Func4: } x_2 := Bx_1$   
→ Lib.Func5(array, bytes);  $\text{Func5: } x_3 := Cx_2 + d$   
→ for (int i = 0; i < 32; i++) {  
→ |→ if (bytes[i] != array[3104 + i]) {  
→ |→ |→ return false;  
→ |→ }  
→ }  
→ return true;  
}
```

is  $x_3 == y$  ?

Pretty easy, eh?

According to static analysis, the input  $x$  should satisfy:

$$C B A x + d = y$$

So, we should be able to get the flag by:

$$flag = A^{-1} B^{-1} C^{-1}(y - d)$$

Pretty easy, eh?

According to static analysis, the input  $x$  should satisfy:

$$C B A x + d = y$$

So, we should be able to get the flag by:

$$flag = A^{-1} B^{-1} C^{-1}(y - d)$$

```
$ sage -python solver.py
```

# Pretty easy, eh?

According to static analysis, the input  $x$  should satisfy:

$$C B A x + d = y$$

So, we should be able to get the flag by:

$$flag = A^{-1} B^{-1} C^{-1}(y - d)$$

```
$ sage -python solver.py
```

```
flag: FAKEFLAGFAKEFLAGFAKEFLAGFAKEFLAG
```

```
$ challenge.exe
```

```
Enter your flag: FAKEFLAGFAKEFLAGFAKEFLAGFAKEFLAG
```

```
Wrong :(
```

# The truth

```
unsafe static Resources()
{
→   IntPtr intPtr = ldftn(Func) - 16;
→   long num = *intPtr;
→   IntPtr intPtr2 = ldftn(Func) - 8;
→   long num2 = *intPtr2;
→   ref long ptr = ldftn(Func) - 16;
→   IntPtr intPtr3 = ldftn(Func) + 5;
→   long num3 = (long)(*(intPtr3 + 1));
→   ptr = *(intPtr3 + (IntPtr)(((int)(*(intPtr3 + 2)) << 3) + 3)) + (num3 << 3);
→   ref long ptr2 = ldftn(Func) - 8;
→   object obj = *(ldftn(Func) - 16);
→   object obj2;
→   for (;;)
→   {
→       obj2 = obj;
→       if (*obj2 == 5)
→       {
→           break;
→       }
→       obj = obj2 + 16;
→   }
→   ptr2 = *(obj2 + 8);
→   long num4 = *(ldftn(Func) - 8);
→   *num4 = 6293447916875450697L;
→   long num5 = num4 + 8L;
→   *num5 = 996842507592L;
```

# About MethodDesc

In .NET, each method has a structure called "MethodDesc" (Method Descriptor).

MethodDesc holds informations like:

- Lower bytes of MethodToken
- Has the method been already JITted?
- Is the method static?
- Method's entry point etc.

There are several types of MethodDesc, but this time we will only focus on the basic one which is used for regular IL methods.

# About MethodDesc

```
0:003> !DumpMT -md 00007fff180e5b00
```

```
(snip)
```

```
MethodDesc Table
```

Entry	MethodDesc	JIT Name
00007fff181f0090	00007fff180e5a98	NONE WCTF2018Rev.Lib.Verify(System.String)
00007fff181f0098	00007fff180e5aa8	NONE WCTF2018Rev.Lib.Func(Int32)
00007fff181f00a0	00007fff180e5ab8	NONE WCTF2018Rev.Lib.Func2()
00007fff181f00a8	00007fff180e5ac8	NONE WCTF2018Rev.Lib.Func3(Byte[], Byte[])
00007fff181f00b0	00007fff180e5ad8	NONE WCTF2018Rev.Lib.Func4(Byte[], Byte[])
00007fff181f00b8	00007fff180e5ae8	NONE WCTF2018Rev.Lib.Func5(Byte[], Byte[])

```
0:003> dq 00007fff180e5a98 l 6
```

00007fff`180e5a98	00280005`20000003	00007fff`181f0090	<- MethodDesc for Verify
00007fff`180e5aa8	00280006`20020004	00007fff`181f0098	<- MethodDesc for Func
00007fff`180e5ab8	00280007`20040005	00007fff`181f00a0	<- MethodDesc for Func2

MethodTokens etc.

Entry point



# Precode to JITted code

PrecodeForFunc1:

```
    call PrecodeFixupThunk  
    pop rsi ; dummy instruction  
    db m_MethodDescChunkIndex_Func1  
    db m_PrcodeChunkIndex_Func1
```

PrecodeForFunc2:

```
    call PrecodeFixupThunk  
    pop rsi ; dummy instruction  
    db m_MethodDescChunkIndex_Func2  
    db m_PrcodeChunkIndex_Func2
```

```
    dq MethodDescBase ; pointer to the first element of  
                        MethodDesc array
```

# Precode to JITted code

PrecodeForFunc1:

```
=> call PrecodeFixupThunk
    pop rsi ; dummy instruction
    db m_MethodDescChunkIndex_Func1
    db m_PrcodeChunkIndex_Func1
    dq MethodDescBase
```

PrecodeFixupThunk:

```
    pop rax
    movzx r10, byte ptr [rax+2] ; m_PrcodeChunkIndex
    movzx r11, byte ptr [rax+1] ; m_MethodDescChunkIndex
    mov rax, qword ptr [rax+r10*8+3] ; MethodDescBase
    lea r10, [rax+r11*8] ; r10=&MethodDesc[r11]
    jmp ThePreStub
```

# Precode to JITted code

PrecodeForFunc1:

```
    call PrecodeFixupThunk
    pop rsi ; dummy instruction    <= rax
    db m_MethodDescChunkIndex_Func1
    db m_PrcodeChunkIndex_Func1
    dq MethodDescBase
```

PrecodeFixupThunk:

```
=> pop rax
    movzx r10, byte ptr [rax+2] ; m_PrcodeChunkIndex
    movzx r11, byte ptr [rax+1] ; m_MethodDescChunkIndex
    mov rax, qword ptr [rax+r10*8+3] ; MethodDescBase
    lea r10, [rax+r11*8] ; r10=&MethodDesc[r11]
    jmp ThePreStub
```

# Precode to JITted code

PrecodeForFunc1:

```
    call PrecodeFixupThunk
    pop rsi ; dummy instruction    <= rax
    db m_MethodDescChunkIndex_Func1
    db m_PrcodeChunkIndex_Func1
    dq MethodDescBase
```

PrecodeFixupThunk:

```
    pop rax
=> movzx r10, byte ptr [rax+2] ; m_PrcodeChunkIndex
    movzx r11, byte ptr [rax+1] ; m_MethodDescChunkIndex
    mov rax, qword ptr [rax+r10*8+3] ; MethodDescBase
    lea r10, [rax+r11*8] ; r10=&MethodDesc[r11]
    jmp ThePreStub
```

# Precode to JITted code

PrecodeForFunc1:

```
    call PrecodeFixupThunk
    pop rsi ; dummy instruction
    db m_MethodDescChunkIndex_Func1
    db m_PrcodeChunkIndex_Func1
    dq MethodDescBase
```

PrecodeFixupThunk:

```
    pop rax
    movzx r10, byte ptr [rax+2] ; m_PrcodeChunkIndex
    movzx r11, byte ptr [rax+1] ; m_MethodDescChunkIndex
=> mov rax, qword ptr [rax+r10*8+3] ; MethodDescBase
    lea r10, [rax+r11*8] ; r10=&MethodDesc[r11]
    jmp ThePreStub
```

# Precode to JITted code

PrecodeForFunc1:

```
    call PrecodeFixupThunk
    pop rsi ; dummy instruction
    db m_MethodDescChunkIndex_Func1
    db m_PrcodeChunkIndex_Func1
    dq MethodDescBase
```

PrecodeFixupThunk:

```
    pop rax
    movzx r10, byte ptr [rax+2] ; m_PrcodeChunkIndex
    movzx r11, byte ptr [rax+1] ; m_MethodDescChunkIndex
    mov rax, qword ptr [rax+r10*8+3] ; MethodDescBase
=> lea r10, [rax+r11*8] ; r10=&MethodDesc[r11]
    jmp ThePreStub
```

# Precode to JITted code

ThePreStub:

```
; save registers to stack  
; (snip)
```

```
; call PreStubWorker
```

```
lea rcx, [rsp+0x68]
```

```
mov rdx, r10 ; MethodDesc
```

```
call PreStubWorker ; do JIT compilation and update MethodDesc
```

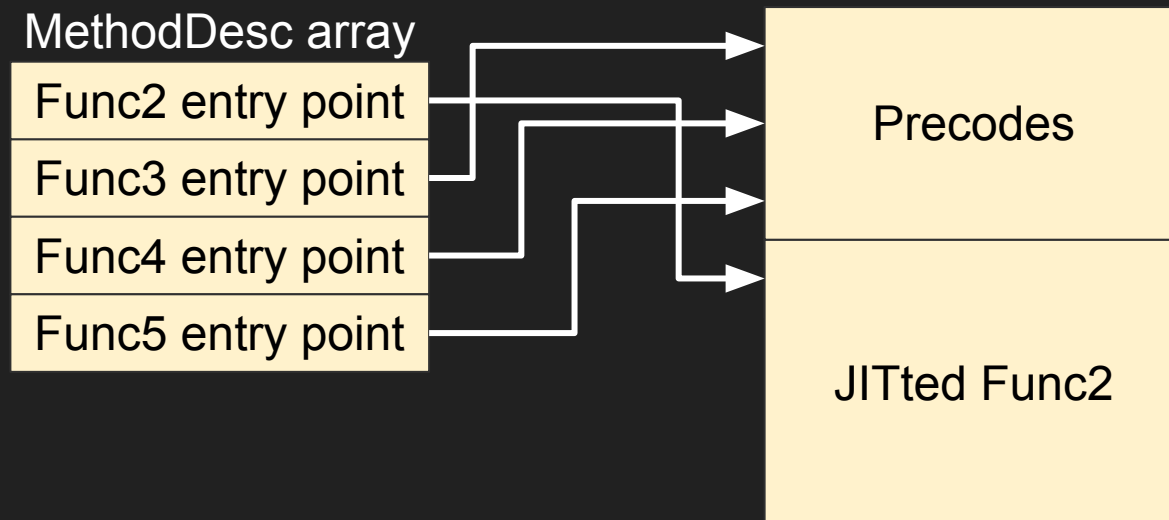
```
; restore registers
```

```
; (snip)
```

```
jmp rax ; jump to compiled code
```

# What Resources.cctor() does

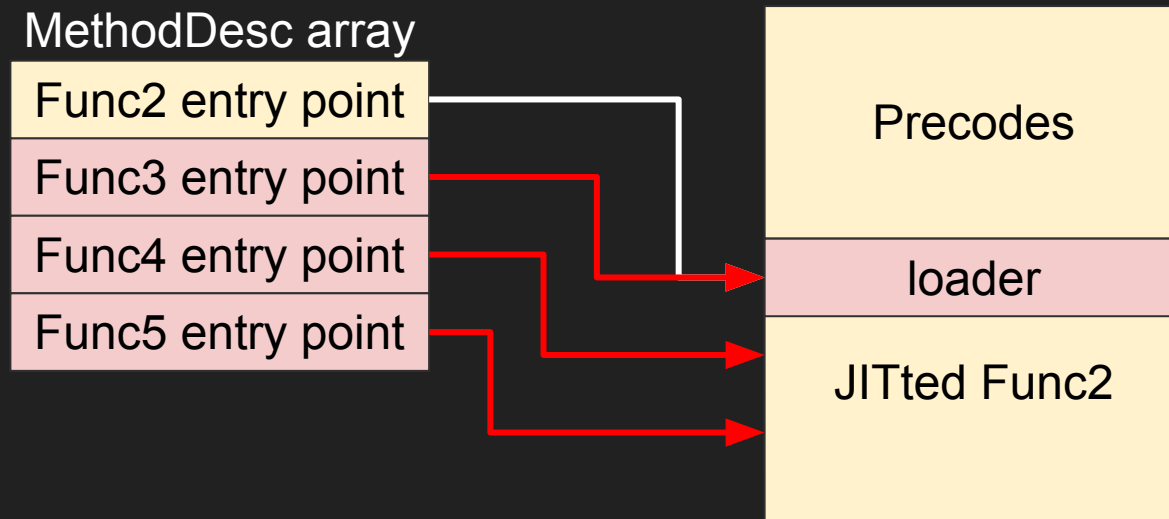
1. Get the entry point of Func1 (which is not JITted yet) by using MSIL ldftn instruction
2. Simulate PrecodeFixupThunk and calculate the address of MethodDesc array
3. Overwrite the beginning of JITted Func2
4. Modify entry points for Func3, 4, and 5





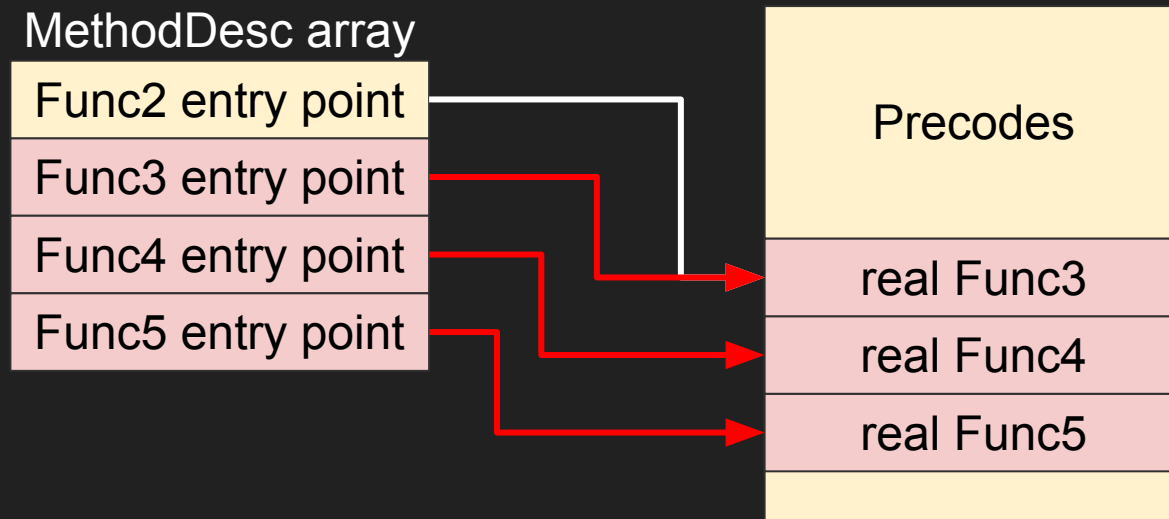
# What Resources.cctor() does

1. Get the entry point of Func1 (which is not JITted yet) by using MSIL Idftn instruction
2. Simulate PrecodeFixupThunk and calculate the address of MethodDesc array
3. Overwrite the beginning of JITted Func2
4. Modify entry points for Func3, 4, and 5



# The real behavior of Func3, 4, and 5

- Func3
  - Extract the real code of Func4 and Func5 from resource
  - $X_1 = AX_0$
- Func4:  $X_2 = B_t^5 X_1 \wedge [0x5a, 0x5a, \dots, 0x5a]$
- Func5:  $10.\text{times}\{ X_3 = C(X_2 \wedge [B_{1,1}, B_{1,2}, \dots, B_{1,32}]) \}$



Fin.